

Enterprise AI Architecture-to-Production

Practitioner Showcase · AI Architecture & Delivery

A methodology and portfolio showcase demonstrating how I engage an AI opportunity, produce evidence-backed architecture decisions, manage gates, and move the initiative toward production, operations, and recertification.

[Live Demo](#) · [GitHub Pages](#) · [15 Scenarios](#) · [330 Artifacts](#)

github.com/danzvla/ai-architecture-to-production

Daniel Mazzini · Principal Architect & Senior TPM

RE & DELIVERY

Architecture-to-Production

Enterprise AI initiatives from initial request through discovery, assessment, architecture, pilot, controlled go-live, operations, and recertification.

the decision gate I manage, and the customer deliverable I produce at every step.

Production Methodology

Methodolog

and grounded in enterprise security, governance, evaluation, and delivery practices · 6 phases · 22 steps.



Intake and Engagement

Initial path before technical design begins.

for

intents

ment

AI Initiative Intake and Engagement Charter

Blank customer-ready template from the MOP deck.

MOP STEP 01 · DISCOVER

01

AI Initiative Intake and Engagement C

One-page charter to confirm sponsor, request, boundary, evidence request, risks and

CUSTOMER / SPONSOR

[Company]: [Sponsor name/title]
Decision coordinator: [name]

STATED REQUEST

[Quote customer request in their own words]
Do not translate into solution yet.

Methodology first, Field MOP second

The showcase is not positioned as another AI tool. It demonstrates a repeatable practitioner route for turning AI demand into defensible architecture, production evidence, and customer-ready deliverables.

01 · Training curriculum

Teaches how enterprise AI decisions are evidenced, defended, and converted into customer artifacts.

02 · Consulting methodology

Structures workshops, discovery, design reviews, readiness gates, and executive recommendations.

03 · Delivery operating model

Turns each step into owners, controls, gates, production evidence, and reusable templates.

04 · Field MOP

Keeps the work procedural enough to repeat: no jump from request to tool selection or production.

Message

Architect can run the engagement, evaluate the opportunity, create the architecture, manage decision gates, and produce results a customer can review and approve.

The method and the evidence are deliberately separated

The original MOP guidance separates the route from the completed examples. The web showcase mirrors that logic through two tabs.

01 - Architecture-to-Production Methodology

Canonical 6-phase, 22-step route. Each step shows the practitioner action, key considerations, gate, and the blank customer-ready template.



02 - Methodology in Practice

Scenario library with 15 realistic enterprise AI engagements. Each scenario has 22 completed artifacts to show how the same method adapts to context.

This prevents confusion: the first tab teaches how I work; the second tab shows the applied evidence.

6 phases, 22 steps, one auditable path to production

Each phase has a defined output and gate so the viewer sees architecture judgment, not random AI content or tool selection.



Technical domains that support the gates

- ✓ AI capability patterns
- ✓ RAG, knowledge & identity
- ✓ Agents, tools & autonomy
- ✓ Security, governance & risk
- ✓ Evaluation & observability
- ✓ Platform, deployment & operations

Engagement discipline before architecture

The early phases show that I can work an AI opportunity before solution design: confirm the ask, align owners, establish evidence, test AI suitability, and prioritize the right use case.

DISCOVER

01

Confirm initiative

02

Frame outcome

03

Map stakeholders

04

Analyze current state

05

Define RACR

ASSESS

06

AI suitability

07

Readiness heatmap

08

Prioritize use case

Capability

I can engage a customer request, clarify business outcome, define decision ownership, expose current state evidence, evaluate AI fit, and produce the artifacts needed to move forward with confidence.

Cross-cutting workstreams run through the lifecycle

Security, data, evaluation, economics, adoption, and operations are not one-time checks. They stay visible across the full MOP.

	DISCOVER	ASSESS	DESIGN	VALIDATE	DELIVER	OPERATE
Business value	●	●	●	●	●	●
Data & knowledge	●	●	●	●	●	●
Security & risk	●	●	●	●	●	●
Responsible AI	●	●	●	●	●	●
Evaluation	●	●	●	●	●	●
FinOps	●	●	●	●	●	●
Architecture decisions	●	●	●	●	●	●
Program governance	●	●	●	●	●	●
Change & adoption	●	●	●	●	●	●
Operations	●	●	●	●	●	●

Field rule: a deep-dive gate does not mean the workstream starts there. Step 13 deepens security/governance; it does not create it. Step 21 operates evaluation/observability; it does not introduce it.

The same lifecycle supports multiple senior AI delivery roles

The role emphasis changes, but the evidence, architecture gates, and production discipline remain consistent.

01 • AI Architect lens

Establish target architecture, ADRs, governance controls, evaluation strategy, and production-readiness gates.

02 • FDE lens

Bridge customer reality to buildable systems through discovery, prototyping, integration, pilot evidence, and rollout support.

03 • Principal SA lens

Align sponsor outcomes, sourcing strategy, economics, risk acceptance, and executive recommendations.


04 • Delivery lens

Run pilot, go-live, service transition, operations, improvement, and recertification as controlled gates.

This is why the showcase is valuable for hiring and consulting: it demonstrates architecture thinking and delivery leadership in the same engagement motion.

Defensible enterprise AI architecture

Design is where architecture patterns, sourcing options, ADRs, logical architecture, governance, evaluation, and implementation topology become explicit.



MOP STEP 11 - DESIGN

AI Architecture Decision Record

Decision record captures that captures context, options, selected decision, rationale, implications and revisit trigger.

DECISION ID / TITLE	CONTEXT	DECISION
ADR-001 [Short Decision Title] Owner: [Name]	[What situation, evidence, requirement or constraint forced this decision?]	[What was decided, stated plainly enough for a sponsor to understand.]

Option	Why considered	Why rejected / assigned	Implication
A	[rationale]	[rationale]	[rationale]
B	[rationale]	[rationale]	[rationale]
C	[rationale]	[rationale]	[rationale]

REVISIT TRIGGER: [] APPROVAL: []

ADRs

MOP STEP 14 - VALIDATE

AI Evaluation and Acceptance Plan

Stakeholders play with golden set, quantify quality, assess threat model and acceptance sign-off.

PRODUCTION/DEVELOPMENT SET	QUALITY	SAFETY	LATENCY / COST
[AI examples] [Source + Location]	[Metric] > [Target] [Validation]	[Metric] > [Target] [Red team scope]	[Metric] < [Target] [Frequency of test]

Test area	Method	Pass threshold	Reviewer	Evidence stored
Functional	[Test type]	[Pass]	[Name]	[Location]
Performance	[Test type]	[Pass]	[Name]	[Location]
Security	[Test type]	[Pass]	[Name]	[Location]

ACCEPTANCE SIGN-OFF: []

Evaluation plan

SCENARIO 01 - TELECOMMUNICATIONS

Security & Governance

AI Threat Model and Governance Control Plan

Threat	Attack / Failure path	Preventive control	Detective control	Owner
Data leakage	Unauthorized access	ACL, feature model	Continuous logs	Security
Resource aging	Not service too broad	Scalable, maintainable tools	Test call audit	Platform
Bad output	Unstable or wrong answer	Guardrails - review	Red team vulnerability	Risk

RISK TIER AND CONTROL POSTURE
Operational Tier 2: no autonomous remediation. Mandatory controls: Human approval, ACL, retrieval, API allowlist, audit trail.

Security design must test AI specific failure modes, not only network and identity controls.

Completed governance artifact

Pilot evidence is not the finish line

Validation, delivery, and operations make productionization explicit: prove, implement, go live carefully, operate, improve, and recertify.

Validate

- 16 Plan controlled pilot
- 17 Evaluate pilot evidence
- 18 Productionization decision gate

Decision: proceed, remediate, or stop.

Deliver

- 19 Production implementation & adoption
- 20 Final readiness & controlled go-live

Decision: release package, canary, rollback, and transition sign-off.

Operate

- 21 Operate & improve
- 22 Recertify / scale / retire

Decision: continue, scale, modify, or retire after material changes.

This is where the showcase distinguishes itself from a typical architecture demo: it keeps production ownership, controls, adoption, runbooks, observability, and recertification visible.

15 scenarios × 22 completed artifacts

The second tab shows how the same method changes with customer context across industries, risks, architecture patterns, systems, and controls.

15

enterprise AI
scenarios

SCENARIO 02 - FINANCIAL SERVICES


Evaluation Criteria
AI Evaluation and Acceptance Plan

Metric	Actual / expected	Target	Status
Model correctness	98.5% rotation accuracy	>98% answer accuracy, >98% rotation accuracy	Pass - conditional
Ground-truth	scenario-specific sample	>95%	Good
Model / feasibility	scenario-linked answers	>95%	Good
Latency / cost	prior telemetry	within approval threshold	Monitor

Next steps / actions Evaluation criteria must be approved before pilot; otherwise the launch decision becomes subjective.

22

completed artifacts
per scenario



330

customer-style
deliverables

2

navigation modes
method + practice

SCENARIO 01 - TELECOMMUNICATIONS

Security & Governance
AI Threat Model and Governance Control Plan

Threat	Attack / Impact path	Prevention control	Detection control	Owner
Data leakage	Unauthorized network	ACL, network isolation	Network access logs	Security
Resource hijacking	Test access too broad	Input validation, auth	Test log audit	Platform
Bad input	Unstable or wrong output	Guardrails - schema	API status telemetry	Risk

RISK TIER AND CONTROL POSTURE
Operational Tier 2: no autonomous remediation. Mandatory controls: Human approval, ACL, retrieval, API allowlist, audit trail.

Next steps / actions Security design must test AI-specific failure modes, not only network and identity controls.

Examples: financial services evaluation, private cloud go-live readiness, and telco security/governance controls.

WHAT THIS DEMONSTRATES

A practitioner capable of running an end-to-end AI opportunity

The showcase proves more than UI design. It demonstrates the ability to engage, assess, architect, govern, validate, deliver, operate, and recertify enterprise AI initiatives.

6

Phases

discover · assess · design · validate · deliver · operate

22

Practitioner steps

actions · considerations · gates · deliverables

330

Completed artifacts

15 scenarios × 22 customer-style deliverables

15

Scenarios

telco · finance · health care · public sector · cloud

100%

Artifact-led

every step has a template or completed deliverable

2-tab

UX model

methodology first, applied scenarios second

TRY IT

Explore the full Architecture-to-Production showcase

Review the methodology, open the templates, switch to Methodology in Practice, and inspect completed scenario deliverables.

✓ • Governed by design

Decision gates make productionization, go-live, operations, and recertification explicit instead of implied.

✓ • Template-driven

Every step is tied to a customer-facing artifact that shows how architecture thinking becomes delivery evidence.

✓ • Scenario-ready

15 worked examples show how the same method adapts to different risk boundaries and operating contexts.

Live demo • <https://danvzla.github.io/ai-architecture-to-production/>

GitHub • github.com/danvzla/ai-architecture-to-production